

Yashwin S

✉ yashwinsblr@gmail.com | in LinkedIn | GitHub

PROFILE

Agentic AI Engineer specializing in designing, developing, and optimizing autonomous workflows, multi-agent orchestration systems, and enterprise RAG solutions on GCP. Hands-on experience building guardrails, tool integrations, and human-in-the-loop workflows to drive operational scalability.

SKILLS

Generative AI: AI Agents, Agent Orchestration, Hermes Agent Framework, Claude Code, LangGraph, LangChain, CrewAI, OpenAI Agents SDK, RAG, Memory Management, Guardrails, Human-in-the-Loop (HITL) Workflows, LlamaIndex, Model Context Protocol (MCP), Prompt Engineering, LLM Evaluation, pgvector, Milvus, Langfuse, OpenTelemetry

Backend: Python, Django, FastAPI, REST APIs, PostgreSQL (RLS), Redis, Supabase, Zitadel Auth, Chargebee, Temporal, Distributed Systems, Microservices Architecture

Cloud/DevOps: GCP, MLOps, Cloudflare Workers AI, Docker, Kubernetes, CI/CD (GitHub Actions, Vercel), SonarQube, Turbopack

Frontend: Next.js, React, TypeScript, Tailwind CSS

EXPERIENCE

Software Engineer - AI

ConnectSecure

Jul 2025 – Present

- **Designed, developed, and optimized an agentic AI Support & Remediation solution** deployed in the ticketing portal to provide security remediation and reduce customer downtime across multiple companies; implemented agent orchestration and tool integration via a Grafana MCP server, Tesseract OCR, and PostgreSQL pgvector hybrid search (+44.5% quality uplift), and applied MLOps telemetry using GLM 5.2 LLM judges and Langfuse tracing.
- **Transformed manual account provisioning processes** into a unified self-service dashboard for B2B tenants and distributors, implementing Django endpoints, transactional subscription workflows, and user sync engines to prevent distributor revenue leakage across cross platform while isolating customer lists via Zitadel JWT claims.
- **Architected and deployed a cloud-based RAG solution on GCP** for parental monitoring, establishing human-in-the-loop (HITL) workflows and custom content guardrails; improved response accuracy by 30% and retrieval precision by 30% using pgvector hybrid search (RRF) while resolving complex timezone and academic calendar edge cases.
- **Designed multimodal ingestion pipelines** (reducing memory by 35% and LLM tokens by 50%) and defined GCP cloud infrastructure (SAD) with GPU/VM latency and billing cost analysis; documented solution designs, technical decisions, and implementation guidelines for scalable maintenance.
- **Optimized Next.js frontend performance** (Turbopack, Next.js 16) and backend query performance for the analytics platform, reducing compile times by 35%, redundant re-renders by 40%, and database overhead by 45% via context state memoization and counts aggregation query caching.

AI Intern

ConnectSecure

Jan 2025 – Jun 2025

- *Sports Medicine Client (Qatar) – Journal RAG:* Built a hybrid-search RAG system over medical PDFs, improving clinician QA performance by 25% across evaluation benchmarks. *Tech:* **Python, JavaScript.**

Research Intern

CDSAML-CCBD, PES University

Jun 2023 – Sep 2023

- Curated **dataset:** 11,800 images, 364 videos (52,568 instances, 6 classes, Indian road conditions). **YOLO v5 pothole detection** (80% accuracy); **MiDaS depth estimation** for automated road monitoring. *Tech:* **Python.**

PUBLICATIONS

Pothole Detection (ACN'23): Built **dataset** and combined **YOLO/EfficientDet** for detection. [Link](#)

Enhancing LLM with CA Knowledge (IEEE ICASET 2025): Implemented **RAG** on **Llama2** with CA Foundation textbook for specialized QA. [Link](#)

CERTIFICATIONS

Claude Code in Action

Anthropic Academy

Certificate

EDUCATION

PES University, Bangalore, India

Dec 2021 – Aug 2025

B.Tech in Computer Science, CGPA: 9.08

- Awarded 6 academic scholarships (MRD, CNR Rao, DAC) for outstanding academic performance.
- Event Lead, AI Coding Club; co-designed smart miner helmet POC (ESP32/Firebase).